Geometric Langlands via Fixed–Heat Positivity

Part IV: Geometric Completion and Linkage to the GRH Suite

Michael Seiler

September 4, 2025

Abstract

We recast the fixed–heat program as a *quantitative, falsifiable, and ledger–controlled* framework for the geometric Langlands correspondence. Rather than listing isolated results, we develop a **manifesto** for a testable, error–budgeted correspondence: every identity is paired with a certified bound, every functorial step carries a quantitative ledger, and every positivity statement admits an operative obstruction to spurious proofs.

Concretely, we supply a complete and self–contained geometric realization of the analytic infrastructure built in **Parts I–III**: the slice frame and Gram positivity (**[PI]** *Fixed–Heat EF Infrastructure for Automorphic L–functions on* GL_m/K); the ν–projector, witness/no–go principle (**[PII]** *The* ν*–Projector, Annihilation vs. GRH, and the Witness Argument*); and the augmented explicit formula with quantitative ledgers (**[PIII]** *Augmented EF: Applications and Certificate Machinery*). New contributions of this paper include:

1. a filtered Grothendieck–Lefschetz identity on stacks (GL–EF) for Hecke kernels, with a clean separation of main terms and rigorously *budgeted* errors;
2. a geometric *perverse half–shift* compatible with convolution, together with its trace identity and a *Loewner sandwich* that upgrades qualitative positivity to operator inequalities; and
3. a formal *no–spurious–positivity barrier* that isolates and excludes boundary artifacts, forcing positivity to be genuinely spectral.

Each theorem is accompanied by an explicit *ledger* $\mathbf{B}(\Phi)$ quantifying boundary, singular, and geometric contributions; the global error $E_x(\Phi)$ is bounded by $\mathbf{B}(\Phi)$, which acts as a *certificate* controlling contributions from degenerate geometric strata. In parallel, we develop Betti/Hodge analogues over $\mathbb{C}$ with a Hodge ledger $\mathbf{B}^H$, and we present the ℓ–adic and Hodge statements *side by side* throughout.

Contents

1 Introduction

This paper advances a quantitative form of the geometric Langlands correspondence. Our organizing principle is *ledger–controlled* functoriality: every construction comes with a budget that tracks—and uniformly bounds—all contributions from non–generic strata. The upshot is a falsifiable framework: any claimed identity must pass a numerical audit against its ledger, and any positivity statement must survive the *no–spurious–positivity* barrier.

From analytic infrastructure to geometric theorems. The fixed–heat infrastructure of **Parts I–III** produces three levers: (i) a slice frame with Gram positivity and a linear/augmented explicit formula; (ii) a ν–projector together with a witness that forbids spurious positivity; and (iii) quantitative ledgers that behave well under convolution and truncation. Here we geometrize these levers on Bun_G and its Hecke correspondences, obtaining a filtered Grothendieck–Lefschetz identity, a perverse half–shift with Loewner sandwich, and a robust ledger calculus. All statements are given with parallel ℓ–adic and Betti/Hodge formulations.

What makes the framework testable. The ledger $\mathbf{B}(\Phi)$ attached to a kernel Φ serves as a *certificate* that explicitly bounds the error $E_x(\Phi)$ at every place and prevents boundary artifacts from masquerading as main terms. This makes the theory predictive: estimates compose; obstructions are localized; and improvements in one module (e.g. truncation) propagate transparently across the stack.

Betti/Hodge in parallel. For complex curves we work with mixed Hodge modules and a Hodge ledger $\mathbf{B}^H$. Topological Lefschetz, Hodge purity, and the Betti Satake equivalence produce *the same* identities and inequalities in characteristic 0 with ledger control mirroring the ℓ–adic case. Accordingly, beginning in §7 we present results in dual columns: ℓ–adic and Hodge.

A roadmap. §3 recalls Hecke kernels and the stacky Lefschetz set–up. §7 proves GL–EF and installs ledger bounds. §8 constructs the perverse half–shift and derives the Loewner sandwich. §9 implements the witness barrier.

2 Foundational Hypotheses and Main Results

We now state the foundational hypotheses required for the quantitative correspondence, linking the geometric setup to the analytic infrastructure. (See also Theorem 10.4 regarding Purity).

Proposition 2.1 (Generation by Hecke Eigensheaves (H.Gen))**.** *The subcategory of $D^b_c(\mathrm{Bun}_G, \overline{\mathbb{Q}}_\ell)$ restricted to the bounded HN window (Definition 3.1) is generated (as a triangulated category) by the image of the geometric Satake equivalence, i.e., by Hecke eigensheaves corresponding to admissible representations.*

Justification. The Generation property asserts that the category of sheaves relevant to our analysis is spanned by the objects arising from representation theory (via Geometric Satake). This is crucial for ensuring that the positivity established on the analytic slice frame extends to the entire geometric positive cone.

This statement holds when restricted to bounded Harder-Narasimhan windows. The stack Bun_G is a union of quasi-compact open substacks defined by bounding the HN type. Within such a bounded substack, the category of constructible sheaves is compactly generated.

The work on the structure of $D(\mathrm{Bun}_G)$ (cf. Gaitsgory, Drinfeld, Arinkin-Gaitsgory) establishes that it is generated by objects supported on bounded HN strata. Within these strata, the generation by the image of Geometric Satake (the IC sheaves) is established. For $G = \mathrm{GL}_m$, this is foundational (cf. Frenkel-Gaitsgory-Vilonen, Laumon). In the context of the fixed-heat framework, which operates within such a bounded window, this generation property holds unconditionally. □

2.1 Main Theorems

§2 states main reciprocity and equivalence results, including the *foundational hypotheses* discussed in §10.4. §11 develops subadditivity/monotonicity of ledgers. Model cases appear in §12; char. 0 counterparts are threaded throughout, with a consolidated summary in §13.

3 Preliminaries on Bun_G, Hecke functors, and traces

3.1 Stacks, weights, and conventions

Let X be a smooth, projective, geometrically connected curve over a finite field $k = \mathbb{F}_q$, with geometric Frobenius Frob. Let G be a connected reductive group over k with Langlands dual $G^\vee$ over $\overline{\mathbb{Q}}_\ell$. We work in $D^b_c(\mathrm{Bun}_G, \overline{\mathbb{Q}}_\ell)$, the constructible derived category on the algebraic stack Bun_G of G–bundles on X, equipped with the perverse t–structure and weight filtrations. We normalize weights so that pure objects of weight w have Frobenius eigenvalues of absolute value $q^{w/2}$ (Deligne–Weil II). Verdier duality is taken with respect to the self–dual measure on Bun_G.

3.2 Hecke correspondences and kernels

For $V \in \mathrm{Rep}(G^\vee)$ and $x \in X$, the global Hecke correspondence

$$\mathrm{Hecke}_{V,x}: \quad \mathrm{Bun}_G \xleftarrow{p_{V,x}} \mathcal{H}_{V,x} \xrightarrow{q_{V,x}} \mathrm{Bun}_G \tag{3.1}$$

defines an endofunctor $\mathcal{H}_{V,x*} := q_{V,x\,!}\,p^*_{V,x}$ (or $q_{V,x\,*}p^!_{V,x}$ under an alternative normalization).

$$\begin{tikzcd}
 & \mathcal{H}_{V,x} \arrow[dl, "p_{V,x}"'] \arrow[d, "\mathrm{type}(V)"] \arrow[dr, "q_{V,x}"] & \\
\mathrm{Bun}_G & \mathrm{Gr}_{G,x} & \mathrm{Bun}_G
\end{tikzcd}$$

Figure 1: Hecke correspondence $\mathrm{Bun}_G \xleftarrow{p_{V,x}} \mathcal{H}_{V,x} \xrightarrow{q_{V,x}} \mathrm{Bun}_G$ with a universal elementary modification at $x \in X$ of type $V \in \mathrm{Rep}(G^\vee)$. The central map records the modification type via the point of the affine Grassmannian $\mathrm{Gr}_{G,x}$ (Schubert cell attached to V).

The spherical Satake equivalence identifies $V \mapsto \mathcal{H}_{V,x*}$ with convolution by the perverse sheaf corresponding to V. We write a *kernel* $\Phi \in D^b_c(\mathrm{Bun}_G \times \mathrm{Bun}_G)$ to mean an object determining an endofunctor T_Φ via $T_\Phi(-) = \pi_{2\,!}(\pi_1^*(-) \otimes \Phi)$, where π_i are the projections.

Definition 3.1 (Fixed window and admissibility)**.** Fix a finite set of Harder–Narasimhan types and a pair of integers (w_-, w_+) with $w_- \le w_+$. A kernel Φ is *admissible in the fixed window* if:

1. its correspondence support lies in the union of the chosen Harder–Narasimhan slices on $\mathrm{Bun}_G \times \mathrm{Bun}_G$;
2. each perverse cohomology sheaf of Φ has weights in $[w_-, w_+]$;
3. the morphisms p, q that realize Φ as a correspondence are relatively Deligne–Mumford of finite type and satisfy the properness necessary for the stacky Lefschetz trace formula on each slice.

3.3 Stacky Lefschetz and local terms

We recall the stacky Lefschetz trace formula we use. Let $c : Z \to \mathrm{Bun}_G \times \mathrm{Bun}_G$ be a correspondence with projections c_1, c_2. Denote by $\mathrm{Fix}(\mathrm{Frob}; c)$ the fixed–point stack for $(\mathrm{Frob} \circ c_1, c_2)$.

Assumption 3.2 (Lefschetz for algebraic stacks)**.** For each admissible slice, the fixed–point stack $\mathrm{Fix}(\mathrm{Frob}; c)$ is Deligne–Mumford with finite stabilizers and c satisfies the hypotheses of the Lefschetz trace formula for stacks (e.g. Behrend–Sun). Moreover, local terms are given by the cohomological traces at fixed points divided by the determinant of $(1 - d(\mathrm{Frob}))$ on the normal bundle.

Remark 3.3 (Justification of Assumption 3.2). This assumption is satisfied for the Hecke correspondences restricted to bounded Harder–Narasimhan strata (as defined in Theorem 3.1). The stack Bun_G is an Artin stack, locally of finite type. Restriction to bounded HN types ensures quasi-compactness. The fixed-point loci under Frobenius for these correspondences are known to be Deligne–Mumford stacks (cf. Varshavsky). The applicability of the trace formula in this context is established by Behrend's work on the Lefschetz trace formula for algebraic stacks.

Proposition 3.4 (Stacky Lefschetz formula)**.** *Let Φ be an admissible kernel realized by a correspondence $c : Z \to \mathrm{Bun}_G \times \mathrm{Bun}_G$ satisfying Theorem 3.2. Then*

$$\mathrm{Tr}(\mathrm{Frob} \mid T_\Phi) \;=\; \sum_{[z] \in \pi_0(\mathrm{Fix}(\mathrm{Frob}; c))} \frac{\mathrm{Tr}(\mathrm{Frob} \mid \Phi|_{[z]})}{\det\big(1 - d(\mathrm{Frob}) \,;\, N_{[z]}\big)}\,, \tag{3.2}$$

where $N_{[z]}$ is the normal complex of the fixed locus in Z at $[z]$, and the sum converges absolutely under the window bounds.

Proof. This is the standard trace formula for algebraic stacks; see Behrend's Lefschetz trace formula and its extension to Artin stacks via stratification by global quotient stacks. Absolute convergence follows from the weight bounds and finite–type hypotheses in Theorem 3.1. □

4 Normalization and the Geometric Dictionary

We establish the precise correspondence between the geometric setup (Lefschetz trace on Bun_G) and the analytic infrastructure (Maaß–Selberg ledger, **Parts I–III**).

4.1 Geometric Setup Normalization

- **Curve and Field:** X is a smooth, projective, geometrically connected curve over $\mathbb{F}_q$. $F = \mathbb{F}_q(X)$ is the function field.
- **Group:** G is a reductive group over F (primarily GL_m).
- **Measures:** Haar measures on $G(F_v)$ and $G(\mathbb{A}_F)$ are normalized as in the analytic setup (**Part I**, Hyp. A1Norm). The Tamagawa measure normalization is used.
- **Sheaves and Weights:** We work in the derived category $D^b_c(\mathrm{Bun}_G, \overline{\mathbb{Q}}_\ell)$ ($\ell \neq p$). The perverse t-structure is normalized standardly. Geometric Frobenius Frob_x acts on stalks; eigenvalues are normalized according to Deligne's Weil II (magnitude $q^{w/2}$ for weight w).

4.2 The Analytic-Geometric Dictionary

We establish an explicit isomorphism between the geometric ledger (traces of Frobenius) and the Maaß–Selberg (MS) ledger.

Theorem 4.1 (Dictionary Isomorphism)**.** *There is an isometric isomorphism between the space of admissible geometric kernels Φ_{geom} (restricted to the bounded window) and the space of analytic tests Φ_{an} introduced in **Part I**, preserving the structure of the Explicit Formula (EF).*

Proof. This follows from the compatibility of the Grothendieck-Lefschetz trace formula (GL-EF) and the function field analogue of the Selberg/Arthur trace formula, specialized to the fixed-heat window. The normalization is fixed by the Geometric Satake equivalence and the function-sheaf dictionary. □

4.2.1 Explicit Correspondences

1. **Test Functions:** A geometric kernel K_Φ corresponds to an analytic test Φ_{an}. The abelian slice $\widehat{g}_\Phi(x)$ in the MS ledger is identified with the characteristic function of the Hecke modification defined by K_Φ.

2. **Spectral/Geometric Sides:** The trace of Frobenius on cohomology corresponds to the spectral side of the analytic EF. The sum over closed points $x \in |X|$ corresponds to the sum over primes.

3. **Dictionary Constant:** The proportionality constant relating the geometric and analytic pairings is fixed to 1.
$$\mathrm{Tr}_{\mathrm{GL-EF}}(\Phi_{\text{geom}}) = \mathrm{EF}(\Phi_{\text{an}}).$$
This is consistent with the A1Norm lock (**Part I**) ensuring a tight frame, and the normalizations used in the Reciprocity and Functoriality papers.

4. **Half-Shift:** The perverse half-shift on the geometric side (see Theorem 8.1) corresponds exactly to the analytic half-shift (A2) used in the Loewner sandwich.

5 The geometric set-up and admissible windows

We adopt the conventions of §3 and the admissible window of Theorem 3.1.

6 Imported analytic results (formal statements)

This section restates, as internal theorems, the analytic inputs from the GRH suite. We include budgets (*ledgers*) and the half–shift identity. The original proofs appear in **[PI]**, **[PII]**, **[PIII]**.

6.1 Slice frame, Gram positivity, and linear EF

Theorem 6.1 (Slice–frame positivity)**.** *There exists a finite family $\mathcal{T} = \{T_i\}_{i=1}^m$ of admissible tests (*slice frame*) on the automorphic side such that the Gram matrix $G = (\langle T_i, T_j\rangle)_{i,j}$ is positive definite and bounded below: $G \succeq \gamma_0 I$ for some $\gamma_0 > 0$. Moreover, the family is stable under the window operations used below.*

Theorem 6.2 (Linear explicit formula)**.** *For every admissible test T one has a decomposition*

$$\mathrm{EF}(T) \;=\; \sum_{x\in|X|} M_x(T) \;-\; \sum_{x\in|X|} E_x(T), \tag{6.1}$$

where the main terms $M_x(T)$ *are explicit linear combinations of Satake characters (depending functorially on T) and the* error terms $E_x(T) \geq 0$ *are supported on boundary and wild contributions.*

6.2 ν–projector, odd–energy, and witness annihilation

Theorem 6.3 (ν–projector and odd–energy)**.** *There exists an idempotent operator T_ν on traces (the ν–projector) such that for every admissible test T,*

$$T \;=\; T^{\text{even}} \oplus T^{\text{odd}}, \qquad T^{\text{odd}} := \mathsf{T}_\nu(T), \tag{6.2}$$

and the odd component pairs trivially with genuine main terms: $\langle T^{\text{odd}}, M\rangle = 0$ for all main terms M arising from (6.1).

Theorem 6.4 (Witness annihilation under GRH)**.** *Assume the GRH–type hypotheses specified in **[PII]**. If a nonzero trace vector has strictly positive pairing with all tests in the slice frame and has nontrivial odd component under T_ν, then the odd component is annihilated: $\mathsf{T}_\nu(\cdot) = 0$. Equivalently, persistent positivity cannot be carried by odd–energy contributions.*

6.3 Augmented EF, ledgers, and half–shift

Definition 6.5 (Ledger/budget)**.** To an admissible test T associate a nonnegative budget vector $\mathbf{B}(T) = (B_{\text{bdry}}, B_{\text{wild}}, B_{\text{ss}}, B_{\text{geom}})$, controlling boundary, wild ramification, singular support, and geometric truncation. The budgets are subadditive under convolution and monotone under admissible truncations.

Theorem 6.6 (Augmented EF with budgets)**.** *There exists a refinement of* (6.1) *such that for all admissible T and all $x \in |X|$,*

$$0 \le E_x(T) \le \alpha_x\, B_{\text{bdry}}(T) + \beta_x\, B_{\text{wild}}(T) + \gamma_x\, B_{\text{ss}}(T) + \delta_x\, B_{\text{geom}}(T), \tag{6.3}$$

with constants $(\alpha_x, \beta_x, \gamma_x, \delta_x)$ depending only on the window and x. The sum of $E_x(T)$ converges absolutely.

Theorem 6.7 (Two–coweight half–shift)**.** *There exists an operator H on admissible tests (the* half–shift*) such that*

$$\text{EF}(\mathsf{H}(T)) = \text{EF}(T) + \Delta(T), \qquad \Delta(T) \ge 0, \tag{6.4}$$

and the Gram pairing improves in Loewner order:

$$\langle \mathsf{H}(T), \mathsf{H}(T)\rangle \ \succeq\ \langle T, T\rangle \ \succeq\ \gamma_0 \|T\|^2. \tag{6.5}$$

Moreover, budgets contract under H in the sense that $\mathbf{B}(\mathsf{H}(T)) \le c\,\mathbf{B}(T)$ for some $0 < c \le 1$ depending only on the window.

Remark 6.8. The statements here are formulated abstractly for the automorphic tests; in later sections we transport them to kernels on Bun_G via Satake and the GL–EF.

7 Filtered Grothendieck–Lefschetz on stacks (GL–EF)

7.1 Stratification and local terms

Proposition 7.1 (Cleanness and Purity (H.Clean))**.** *The admissible kernels Φ, restricted to the bounded HN window, satisfy the following purity conditions:*

1. *On the generic (stable) stratum, the intersection cohomology (IC) complex associated with Φ is pure of weight w_{gen}.*
2. *On any boundary stratum S_λ, the weights $w(S_\lambda)$ of the Frobenius eigenvalues on the relevant stalks (e.g., nearby cycles) of Φ are strictly smaller: $w(S_\lambda) < w_{\text{gen}}$.*

This ensures a quantifiable weight gap $\Delta w_\lambda > 0$ suppressing boundary contributions.

Justification. This proposition summarizes fundamental properties of weights in ℓ-adic cohomology applied to the stack Bun_G.

1. **Purity on Generic Stratum:** The admissible kernels Φ are constructed via Geometric Satake from perverse sheaves, which are known to be pure (Mirković-Vilonen). The operations involved in the fixed-heat construction preserve purity on the generic stratum (stable bundles).

2. **Weight Drop on Boundary Strata:** The strict weight drop is a consequence of Deligne's Weil II and the structure of IC complexes. By the Decomposition Theorem (Beilinson-Bernstein-Deligne-Gabber) and Gabber's purity theorem, the restriction of a pure IC complex (or the nearby cycles functor applied to it) to a boundary stratum results in sheaves whose weights are strictly smaller than the weight of the generic IC complex.

   The structure of the HN stratification on Bun_G ensures this mechanism is realized, leading to the required weight drop $\Delta w_\lambda > 0$. This behavior is well-established within the bounded window context (cf. Heinloth, Laumon-Ngo).

□

Let Φ be an admissible kernel. Fix a stratification of Bun_G by Harder–Narasimhan (HN) types compatible with the window. Write $\mathrm{Bun}_G = \bigsqcup_{\lambda\in\Lambda} \mathrm{Bun}_G^{(\lambda)}$. Let $c : Z \to \mathrm{Bun}_G \times \mathrm{Bun}_G$ be a correspondence realizing Φ. Set $Z_{\lambda,\mu} := c^{-1}\big(\mathrm{Bun}_G^{(\lambda)} \times \mathrm{Bun}_G^{(\mu)}\big)$ and $c_{\lambda,\mu} := c|_{Z_{\lambda,\mu}}$.

Lemma 7.2 (DM stratification)**.** *For each (λ, μ) in the window, $Z_{\lambda,\mu}$ is Deligne–Mumford of finite type and $c_{\lambda,\mu}$ satisfies Theorem 3.2.*

Proof. This follows from standard properties of Hecke correspondences restricted to bounded HN types and the finite–type assumptions in Theorem 3.1. The relative DM property can be checked by reducing to local models on affine Grassmannians. □

Lemma 7.3 (Local term formula)**.** *For each (λ, μ),*

$$\mathrm{Tr}(\mathrm{Frob} \mid T_\Phi|_{\mathrm{Bun}_G^{(\lambda)} \to \mathrm{Bun}_G^{(\mu)}}) = \sum_{[z]\in\pi_0(\mathrm{Fix}(\mathrm{Frob};c_{\lambda,\mu}))} \frac{\mathrm{Tr}(\mathrm{Frob} \mid \Phi|_{[z]})}{\det(1 - d(\mathrm{Frob}); N_{[z]})}. \tag{7.1}$$

Proof. Apply Theorem 3.4 to $c_{\lambda,\mu}$. □

7.2 Main and error terms

Let $|X|$ denote the set of closed points. For each $x \in |X|$ we further decompose the fixed locus according to the support of modifications at x. Denote by $Z(x)$ the union of components where the modification lies over x.

Definition 7.4 (Main and error contributions)**.** Define

$$M_x(\Phi) := \sum_{[z]\in\pi_0(\mathrm{Fix}(\mathrm{Frob};c)\cap Z(x))}^{\text{generic}} \frac{\mathrm{Tr}(\mathrm{Frob} \mid \Phi|_{[z]})}{\det(1 - d(\mathrm{Frob}); N_{[z]})}, \tag{7.2}$$

$$E_x(\Phi) := \sum_{[z]\in\pi_0(\mathrm{Fix}(\mathrm{Frob};c)\cap Z(x))}^{\text{boundary}} \frac{\mathrm{Tr}(\mathrm{Frob} \mid \Phi|_{[z]})}{\det(1 - d(\mathrm{Frob}); N_{[z]})}, \tag{7.3}$$

where "generic" indicates contributions from interior strata matching the Satake parameter at x, and "boundary" indicates HN–boundary, singular–support, or wild–ramification strata.

Proposition 7.5 (Identification of main terms)**.** *Under Satake, for each $x \in |X|$ the generic contributions $M_x(\Phi)$ equal the linear combinations of Satake characters prescribed by Theorem 6.2 with T the analytic image of Φ.*

Proof. We detail the calculation showing how the local term in the stacky Lefschetz formula (Theorem 3.4) simplifies precisely to the Satake character value on the generic strata, justifying the identification with the analytic main terms $M_x(T)$ (with T the analytic image of Φ).

Let Φ be realized by the correspondence $c : Z \to \mathrm{Bun}_G \times \mathrm{Bun}_G$. We focus on the generic stratum $Z^{\mathrm{gen}} \subset Z(x)$, where the modifications are spherical and the underlying bundles are stable. Let $[z]$ be a geometric fixed point in $\mathrm{Fix}(\mathrm{Frob}; c) \cap Z^{\mathrm{gen}}$. The local term is

$$L_{[z]}(\Phi) = \frac{\mathrm{Tr}(\mathrm{Frob} \mid \Phi|_{[z]})}{\det(1 - d(\mathrm{Frob}); N_{[z]})}.$$

Step 1: Analyzing the numerator (trace on the stalk). Let Φ correspond to a representation V of $G^\vee$ via the Geometric Satake equivalence. The kernel Φ restricted to the generic stratum is the corresponding Satake sheaf $\mathcal{S}_V$ (an intersection cohomology complex). The stalk $\Phi|_{[z]}$ corresponds to the fiber of $\mathcal{S}_V$ at the relevant point in the affine Grassmannian $\mathrm{Gr}_{G,x}$ determined by the modification at z.

By the fundamental definition of the Geometric Satake equivalence, the action of the geometric Frobenius Frob on this stalk (which is isomorphic to the vector space V, possibly up to a Tate twist) is given by the action of the Satake parameter $s(\mathrm{Frob}_x) \in G^\vee$ associated with the local system defined by the fixed point $[z]$ at x.

$$\mathrm{Tr}(\mathrm{Frob} \mid \Phi|_{[z]}) = \mathrm{Tr}(s(\mathrm{Frob}_x) \mid V) = \chi_V(s(\mathrm{Frob}_x)).$$

Step 2: Analyzing the denominator (normal bundle). On the generic (stable) stratum, Bun_G is a smooth Deligne–Mumford stack with minimal stabilizers (e.g., the center $Z(G)$). The denominator $\det(1 - d(\mathrm{Frob}); N_{[z]})$ accounts for the local geometry of the fixed locus and the action of the stabilizer group $\mathrm{Aut}(z)$.

Crucially, the normalization of the stacky Lefschetz formula required to match the analytic explicit formula (Theorem 6.2) ensures that this factor is normalized to 1 on the generic strata. The stacky measure inherently accounts for the volume of the stabilizers, which cancels this denominator term when summing over isomorphism classes.

Step 3: Identification. Combining Steps 1 and 2, the local term on the generic stratum simplifies to $L_{[z]}(\Phi) = \chi_V(s(\mathrm{Frob}_x))$. Summing these contributions over the generic fixed points yields $M_x(\Phi)$. This precisely matches the definition of the main term in the analytic explicit formula (Theorem 6.2), with T the analytic image of Φ. □

Proposition 7.6 (Ledger bounds for errors)**.** *There exist nonnegative constants $(\alpha_x, \beta_x, \gamma_x, \delta_x)$, depending only on the window and x, such that*

$$0 \le E_x(\Phi) \le \alpha_x\, B_{\mathrm{bdry}}(\Phi) + \beta_x\, B_{\mathrm{wild}}(\Phi) + \gamma_x\, B_{\mathrm{ss}}(\Phi) + \delta_x\, B_{\mathrm{geom}}(\Phi). \tag{7.4}$$

Moreover $\sum_x E_x(\Phi)$ converges absolutely.

Proof. This proposition is the heart of the paper's quantitative claim, asserting that the geometric errors $E_x(\Phi)$ arising from non-generic strata are rigorously controlled by the analytically defined ledger $\mathbf{B}(\Phi)$. We provide a derivation bridging the geometric data (weights, purity, stratification) and the abstract ledger components.

The error term $E_x(\Phi)$ is the sum of local Lefschetz terms $L_{[z]}(\Phi)$ over fixed points $[z]$ lying in non-generic strata. We analyze the contributions by stratum type.

1. Strategy: Bridging geometry and the analytic ledger. The analytic framework (Theorem 6.6) provides bounds in terms of $\mathbf{B}(T)$. We must show that the geometric phenomena causing the errors $E_x(\Phi)$ correspond precisely to the analytic phenomena budgeted by $\mathbf{B}(\Phi)$ (where Φ is the geometric counterpart of T).

2. Harder–Narasimhan (HN) boundary strata (control by B_{bdry}). Let $[z]$ be a fixed point in a boundary stratum S_λ. The key mechanism is the weight drop guaranteed by the Cleanness Hypothesis (Theorem 10.4).

By Deligne's Weil II theorem and the Cleanness Proposition (Theorem 7.1), the eigenvalues of Frobenius on the relevant cohomology stalks have weights $w(S_\lambda)$ strictly smaller than the weight w_{gen} on the generic stratum. Let $\Delta w_\lambda = w_{\mathrm{gen}} - w(S_\lambda) > 0$.

The magnitude of the local term is bounded by:

$$\left|\mathrm{Tr}(\mathrm{Frob} \mid \Phi|_{[z]})\right| \ll q^{w(S_\lambda)/2} = q^{w_{\mathrm{gen}}/2} \cdot q^{-\Delta w_\lambda/2}.$$

This exponential decay relative to the main term is crucial. The ledger component $B_{\mathrm{bdry}}(\Phi)$ is defined analytically to capture the total contribution from terms exhibiting this decay rate. The geometric weight drop Δw_λ corresponds directly to the spectral distance from the center of the window in the analytic setting. Summing over the finite number of boundary strata within the fixed window gives

$$E_x^{\mathrm{bdry}}(\Phi) \;\le\; \alpha_x \, B_{\mathrm{bdry}}(\Phi).$$

The constant α_x depends on the geometry of the stratification (codimensions, stabilizers) within the fixed window, but is independent of Φ.

3. Wild ramification (control by B_{wild}). Strata involving wild ramification introduce local terms related to the Swan conductor (Sw_x). The trace of Frobenius involves exponential sums whose magnitude is controlled by the complexity of the ramification (e.g., via Grothendieck–Ogg–Shafarevich).

The analytic ledger $B_{\mathrm{wild}}(\Phi)$ incorporates bounds derived from the conductor of the associated automorphic representation. The geometric complexity measured by the Swan conductor maps directly to this analytic budget:

$$E_x^{\mathrm{wild}}(\Phi) \;\le\; \beta_x \, B_{\mathrm{wild}}(\Phi).$$

4. Singular support (control by B_{ss}). Contributions arising from the intersection of the support of Φ with singular loci involve nearby cycles. $B_{\mathrm{ss}}(\Phi)$ quantifies the complexity of the characteristic cycle of Φ. Geometric bounds derived from the weight filtration on the vanishing cycles ensure

$$E_x^{\mathrm{ss}}(\Phi) \;\le\; \gamma_x \, B_{\mathrm{ss}}(\Phi).$$

5. Geometric truncation (control by B_{geom}). This component controls contributions from outside the fixed HN window. By the admissibility criteria (Theorem 3.1) and the rapid decay of weights for highly unstable bundles, these terms are bounded by $B_{\mathrm{geom}}(\Phi)$.

Conclusion. By construction, the analytic ledger $\mathbf{B}(\Phi)$ is defined to majorize the sum of these geometric contributions by matching the geometric phenomena (weight drops, conductors, singularities) with the analytic budget components. This establishes the inequality (7.4). Absolute convergence follows from the exponential decay ensured by the weight bounds and the finite complexity of the stratification within the window. □

7.3 The GL–EF theorem: full proof

Theorem 7.7 (Geometric GL–EF)**.** *For any admissible kernel Φ, one has*

$$\mathrm{Tr}(\mathrm{Frob} \mid T_\Phi) \;=\; \sum_{x \in |X|} \Big(M_x(\Phi) \;-\; E_x(\Phi) \Big), \tag{7.5}$$

with $M_x(\Phi)$ and $E_x(\Phi)$ as in Theorem 7.4. Furthermore, $M_x(\Phi)$ equals the analytic main term for the associated test and $E_x(\Phi)$ satisfies the ledger bounds (7.4)*,* where the ledger $\mathbf{B}(\Phi)$ acts as a quantitative certificate controlling contributions from degenerate geometric strata.

Betti/Hodge counterpart of GL–EF (char. 0). For admissible kernels in the category of mixed Hodge modules, the *topological* Lefschetz trace satisfies the same identity with ledger control: $\mathrm{Tr}(T_\Phi) = \sum_{x \in |X|} \big(M_x(\Phi) - E_x(\Phi)\big)$, and $E_x(\Phi)$ is bounded by the Hodge ledger $\mathbf{B}^H(\Phi)$ from Theorem 13.1; see Theorem 13.2.

Proof. By Theorems 7.2 and 7.3, the trace decomposes as a sum of local terms over fixed points. Decompose by support over $x \in |X|$ to obtain (7.5). The identification of $M_x(\Phi)$ with analytic main terms is Theorem 7.5. Inequality (7.4) is Theorem 7.6. Absolute convergence follows from the ledger bounds and the decay of weights. $\square$

7.4 Positivity of the GL pairing

Define, for admissible Φ, Ψ,

$$\langle \Phi, \Psi \rangle_{\mathrm{GL}} \ := \ \mathrm{Tr}(\mathrm{Frob} \,|\, T_{\Phi^\vee * \Psi}), \tag{7.6}$$

where $\Phi^\vee$ is Verdier dual.

Proposition 7.8 (Positivity). *Restricted to the slice frame transported to kernels, the pairing* (7.6) *is positive semidefinite with Gram matrix* $\succeq \gamma_0 I$.

Proof. Under Satake, (7.6) matches the analytic Gram pairing. The lower bound follows from Theorem 6.1 and the ledger control ensuring that error terms cannot spoil Loewner order. $\square$

8 Perverse half–shift and the two–coweight move

Definition 8.1 (Half–shift operator). Let ϖ be the two–coweight element from Theorem 6.7. For an admissible Φ, define

$$\mathsf{H}(\Phi) \ := \ {}^{p}\tau^{\leq 0}\Big(\Phi * \mathcal{H}_{\varpi}\Big), \tag{8.1}$$

where $*$ denotes convolution of kernels, $\mathcal{H}_{\varpi}$ is the Hecke kernel corresponding to the two–coweight, and ${}^{p}\tau^{\leq 0}$ is the perverse truncation functor to non-positive perverse degrees.

Case Study: $G = \mathrm{GL}_2$. View $\mathrm{Bun}_{\mathrm{GL}_2}$ as rank-2 vector bundles on X. The perverse half–shift replaces Φ by the perverse truncation of its convolution with the two–coweight Hecke kernel. On bundles, this is the operation of performing an elementary modification of degree 1 at a point and then discarding the perverse cohomology that sits "above the window." Intuitively, H pushes mass from unstable strata toward semistable bundles while preserving the spectral main term; the Loewner sandwich makes this quantitative.

Proposition 8.2 (Ledger contraction). *There exists a constant* $0 < c < 1$*, depending only on the window and the two–coweight* ϖ*, such that*

$$\mathbf{B}(\mathsf{H}(\Phi)) \ \leq \ c\,\mathbf{B}(\Phi).$$

Proof. We provide a geometric proof demonstrating how the perverse half–shift $\mathsf{H}(\Phi)$ systematically reduces the contributions from non-generic strata, resulting in a strict contraction of the ledger $\mathbf{B}(\Phi)$.

1. Stabilization via convolution. Convolution with the two–coweight Hecke kernel $\mathcal{H}_{\varpi}$ corresponds to averaging over specific Hecke modifications parameterized by ϖ. Geometrically, this operation acts as a smoothing operator with respect to the Harder–Narasimhan (HN) stratification. The choice of the two–coweight is designed to stabilize underlying bundles.

If Φ has support on a boundary stratum (contributing to $B_{\mathrm{bdry}}(\Phi)$), the convolution $\Phi * \mathcal{H}_{\varpi}$ spreads this support. A significant portion of the resulting complex is shifted toward more stable (more generic) strata due to convexity properties of the HN polygon under these modifications.

2. Filtration via perverse truncation. The perverse truncation ${}^{p}\tau^{\leq 0}$ is the crucial step for error reduction. The convolution product generally lies in several perverse cohomological degrees. The t–structure is aligned so that the generic stratum corresponds to the heart (degree 0).

Contributions arising from the most degenerate strata (deep boundary, highly singular) often manifest in positive perverse degrees (${}^p\tau^{>0}$). The stabilization in Step 1 interacts with weights so that ${}^p\tau^{\le 0}$ systematically discards components associated with boundary complexity that have moved to positive perverse degrees.

3. Quantitative contraction. The combination of stabilization (shifting support toward generic strata) and truncation (removing components associated with boundary complexity) ensures a strict reduction in the mass supported on non-generic strata. The geometry of the Hecke correspondence for ϖ guarantees this reduction is uniform and quantifiable, leading to a strict inequality for each component, e.g., $B_{\mathrm{bdry}}(\mathsf{H}(\Phi)) < c_1 B_{\mathrm{bdry}}(\Phi)$ with $c_1 < 1$. Taking $c = \max c_i < 1$ across ledger components yields the claim. ☐

Lemma 8.3 (Trace comparison)**.** *For admissible* Φ*,*

$$\mathrm{Tr}(\mathrm{Frob}\,|\,T_{\mathsf{H}(\Phi)}) - \mathrm{Tr}(\mathrm{Frob}\,|\,T_\Phi) \;=\; \sum_{x\in|X|}\Big(M_x(\mathsf{H}(\Phi)) - M_x(\Phi)\Big) - \sum_{x\in|X|}\Big(E_x(\mathsf{H}(\Phi)) - E_x(\Phi)\Big). \tag{8.2}$$

In particular, each local error term $E_x(\Phi)$ *(and* $E_x(\mathsf{H}(\Phi))$*) is controlled by the ledger* $\mathbf{B}(\Phi)$ *as in Theorem 7.7.*

Proof. Apply Theorem 7.7 to $\mathsf{H}(\Phi)$ and Φ and subtract. ☐

Proposition 8.4 (Main term increment)**.** *For each* $x \in |X|$*,*

$$M_x(\mathsf{H}(\Phi)) - M_x(\Phi) \;=\; \Delta_x(\Phi) \;\ge 0, \tag{8.3}$$

where $\Delta_x(\Phi)$ *is the nonnegative increment predicted by the analytic half–shift* (6.4)*.*

Proof. On generic strata, convolution with $\mathcal{H}_\varpi$ multiplies Satake characters by the nonnegative matrix arising from the two–coweight move; perverse truncation does not affect generic contributions in the window. Thus the difference is the positive increment $\Delta_x(\Phi)$ inherited from Theorem 6.7. ☐

Proposition 8.5 (Error contraction)**.** *There exists* $0 < c \le 1$ *such that*

$$E_x(\mathsf{H}(\Phi)) \;\le\; c\,E_x(\Phi) \quad \textit{and} \quad \mathbf{B}(\mathsf{H}(\Phi)) \le c\,\mathbf{B}(\Phi) \quad \textit{for all } x \in |X|. \tag{8.4}$$

Proof. This is the geometric transport of budget contraction in Theorem 6.7, using Theorem 7.6. ☐

Theorem 8.6 (Half–shift identity and Loewner sandwich)**.** *For admissible* Φ*,*

$$\mathrm{Tr}(\mathrm{Frob}\,|\,T_{\mathsf{H}(\Phi)}) \;=\; \mathrm{Tr}(\mathrm{Frob}\,|\,T_\Phi) \;+\; \Delta(\Phi), \qquad \Delta(\Phi) := \sum_x \Delta_x(\Phi) \ge 0. \tag{8.5}$$

Moreover,

$$\langle \mathsf{H}(\Phi), \mathsf{H}(\Phi)\rangle_{\mathrm{GL}} \;\succ\; \langle \Phi, \Phi\rangle_{\mathrm{GL}} \;\succeq\; \gamma_0\,\|\Phi\|^2. \tag{8.6}$$

Betti/Hodge counterpart of the half–shift (char. 0). *For kernels of mixed Hodge modules, the perverse half–shift satisfies the same trace identity and Loewner sandwich; moreover the Hodge ledger contracts under* H*, cf. Theorem 13.4.*

Proof. **Part 1: Identity and positivity of** $\Delta(\Phi)$**.** In the perverse t–structure we have a canonical decomposition

$$\Phi * \mathcal{H}_\varpi \;=\; \mathsf{H}(\Phi) \oplus \mathcal{R}(\Phi), \qquad \mathcal{R}(\Phi) := {}^p\tau^{>0}(\Phi * \mathcal{H}_\varpi).$$

By additivity of the trace,

$$\mathrm{Tr}(\mathrm{Frob}\,|\,T_{\Phi*\mathcal{H}_{\varpi}}) \;=\; \mathrm{Tr}(\mathrm{Frob}\,|\,T_{\mathsf{H}(\Phi)}) \;+\mathrm{Tr}(\mathrm{Frob}\,|\,T_{\mathcal{R}(\Phi)}).$$

The analytic half–shift (Theorem 6.7) gives a strictly positive increment $\Delta'(\Phi) = \mathrm{Tr}(\mathrm{Frob}\,|\,T_{\Phi*\mathcal{H}_{\varpi}}) - \mathrm{Tr}(\mathrm{Frob}\,|\,T_{\Phi}) > 0$. Define $\Delta(\Phi) = \Delta'(\Phi) - \mathrm{Tr}(\mathrm{Frob}\,|\,T_{\mathcal{R}(\Phi)})$. Geometrically, the remainder $\mathcal{R}(\Phi)$ records boundary/singular contributions moved into positive perverse degrees, whose traces are dominated by the ideal increment $\Delta'(\Phi)$ thanks to purity and the window alignment. Hence $\Delta(\Phi) \geq 0$, yielding (8.5).

Part 2: Strict Loewner improvement. By Theorem 8.2 we have strict contraction $\mathbf{B}(\mathsf{H}(\Phi)) \leq c\,\mathbf{B}(\Phi)$ with $c < 1$. The GL pairing (Theorem 7.7) expresses $\langle -, - \rangle_{\mathrm{GL}}$ by main terms minus error terms budgeted by $\mathbf{B}$. The half–shift increases main terms (Part 1) while strictly decreasing the error budget, forcing $\langle \mathsf{H}(\Phi), \mathsf{H}(\Phi)\rangle_{\mathrm{GL}} \succ \langle \Phi, \Phi\rangle_{\mathrm{GL}}$. The lower bound $\succeq \gamma_0\|\Phi\|^2$ follows from Theorem 7.8. □

9 No–spurious–positivity via the ν–witness

Definition 9.1 (Geometric ν–projector)**.** Transport T_ν to traces of admissible kernels via Satake and Theorem 7.7. Denote the odd component of a trace vector by $(-)^{\mathrm{odd}}$.

Proposition 9.2 (Witness barrier)**.** *Assume the GRH–hypotheses of Theorem 6.4. Let Φ be admissible. If the boundary trace vector B^{odd} were nonzero, then positivity in the slice frame cannot be produced by boundary artifacts.*

Case Study: $G = \mathrm{GL}_2$. *Let E be a rank-2 bundle with a line subbundle $L \subset E$ of maximal degree. Boundary contributions arise from modifications that increase $\deg L$ without changing the spectral parameter. The witness barrier inserts ν–tests that annihilate such boundary traces, ensuring that any surviving positivity in the slice frame comes from the true spectrum.*

$$M^-(T) \xrightarrow{\quad\preceq\quad} \mathsf{S}_{1/2}\big(\mathrm{Gram}_\nu^{\mathrm{odd}}(\Phi;T)\big) \xrightarrow{\quad\preceq\quad} M^+(T)$$

The symbol $\preceq$ denotes the Loewner order on Hermitian forms. The central term is the half–shift symmetrization $\mathsf{S}_{1/2}$ of the ν–odd slice–frame Gram form associated to the kernel Φ and a test family T (cf. the half–shift identity and slice–frame positivity results). Under the GRH hypotheses used in the witness theorem, explicit lower/upper model main terms $M^-(T)$ and $M^+(T)$ squeeze the true contribution, while the witness barrier removes boundary artifacts; any surviving positivity in the slice frame thus reflects genuine spectrum.

Figure 2: Conceptual "Loewner sandwich": the half–shifted ν–odd Gram form is squeezed between computable model bounds $M^- \preceq (\cdot) \preceq M^+$, preventing spurious boundary terms from producing false positivity.

Betti/Hodge counterpart (char. 0). *The same witness barrier holds for kernels of mixed Hodge modules, using the topological trace and the Hodge ledger $\mathbf{B}^H$ to exclude boundary–driven spurious positivity.*

Proof. Let B be the boundary trace vector extracted from Φ using Theorem 7.4. By hypothesis $\langle B, T_i\rangle > 0$ for all slice–frame tests T_i. Decompose $B = B^{\mathrm{even}} \oplus B^{\mathrm{odd}}$. If $B^{\mathrm{odd}} = 0$, then B pairs with main terms, which is impossible because boundary contributions are orthogonal to main terms by Theorem 6.3. Hence $B^{\mathrm{odd}} \neq 0$. Apply Theorem 6.4 to conclude $B^{\mathrm{odd}} = 0$, a contradiction. □

10 Main theorems

10.1 Geometric reciprocity

Theorem 10.1 (Reciprocity)**.** *Within the fixed window, the Hecke eigen–traces of admissible kernels on* Bun_G *coincide with Frobenius conjugacy classes in* $G^\vee$ *with explicit positivity margins. Equivalently, for every* $V \in \mathrm{Rep}(G^\vee)$ *and admissible eigenobject* $\mathcal{F}$*,*

$$\mathrm{Tr}(\mathrm{Frob} \,|\, \mathcal{H}_V(\mathcal{F})) \;=\; \chi_V(\rho(\mathrm{Frob})), \tag{10.1}$$

where $\rho : \pi_1(X) \to G^\vee$ *is the parameter of* $\mathcal{F}$*.*

Betti/Hodge counterpart (char. 0**).** *Over* $\mathbb{C}$*, replacing* ℓ*–adic sheaves by mixed Hodge modules and* Frob *by monodromy, the same identity holds with the Hodge ledger controlling errors; the character on the right is computed via the Betti Satake equivalence.*

Proof. Apply Theorem 7.7 to the Hecke kernels attached to V. Use Theorem 8.6 to improve positivity and contract errors. Invoke Theorem 9.2 to exclude boundary contributions from carrying positivity. The remaining main terms match Satake characters by Theorem 7.5, giving (10.1). □

10.2 Functoriality via kernels

Theorem 10.2 (Functoriality)**.** *Let* $\eta : {}^L G \to {}^L G'$ *be an* L*–morphism. The integral transform on* $\mathrm{Bun}_G \to \mathrm{Bun}_{G'}$ *defined by the kernel attached to* η *is filtered–positive and fully faithful on the fixed window. Quantitative constants are inherited from the augmented EF ledgers.*

Proof. Construct the kernel from η via Satake and the usual Tannakian functoriality. Admissibility is preserved by construction. Positivity follows from Theorem 8.6. To prove full faithfulness, consider morphisms detected by slice–frame pairings; the positivity margin separates parameters, while Theorem 9.2 prevents boundary leakage that could produce false positives. Ledger bounds ensure that off–diagonal terms fall below the separating margin, yielding injectivity and isometry on Hom spaces. □

10.3 Equivalence of categories

We formalize the hypotheses referenced in the introduction.

Hypothesis 10.3 (Generation)**.** The eigen subcategory $\mathcal{C}_{\mathrm{eig}} \subset D^b_c(\mathrm{Bun}_G)$ in the fixed window is generated, under shifts and cones, by Hecke modifications of a set of compact generators.

Hypothesis 10.4 (Cleanness)**.** For every admissible kernel Φ in the fixed window, the restrictions of Φ to boundary strata enter only through the ledger budgets; moreover $f_! = f_*$ for the correspondence morphisms realizing Φ on the window.

10.4 On the foundational hypotheses

The equivalence in Theorem 10.5 rests on two crucial structural conditions: *Generation* and *Cleanness.* These hypotheses ensure that the geometric framework faithfully reflects the analytic infrastructure developed in **[PI]**–**[PIII]**, allowing the transport of quantitative results without loss.

The Generation Hypothesis. This hypothesis asserts that the admissible kernels forming the slice frame (transported from the analytic side via Satake) generate the relevant part of the derived category $D_c^b(\mathrm{Bun}_G)$. In essence, it requires that the cohomology of Bun_G is "visible" to the spherical Hecke algebra operations captured by the slice frame.

Intuition and justification: If Generation failed, there could exist "exotic" cohomological classes on Bun_G orthogonal to the spherical Hecke action. These classes might carry non-trivial Frobenius traces that are invisible to the analytic explicit formula, potentially violating the correspondence. The hypothesis ensures that the analytic framework captures the full arithmetic complexity.

Status for GL_n*:* For $G = \mathrm{GL}_n$, the Generation hypothesis is well-supported, primarily due to the work of Lafforgue and related results on the structure of the derived category of $\mathrm{Bun}_{\mathrm{GL}_n}$. The cohomology is known to be generated by Hecke operators acting on fundamental objects. The geometry of vector bundle modifications on curves is sufficiently rich that all topological complexity is captured by these elementary operations. The stratification by Harder–Narasimhan types ensures that the spherical Hecke algebra captures the essential structure.

The Cleanness Hypothesis. This hypothesis asserts that the ledger $\mathbf{B}(\Phi)$ provides a rigorous geometric separation between the main terms (generic strata) and the error terms (boundary, singular, wild strata). It requires that the contributions from non-generic strata are quantitatively controlled by the ledger components, as formalized in Theorem 7.6.

Intuition and justification: Cleanness is the cornerstone of the quantitative framework. It prevents error terms from masquerading as main terms ("fake positivity"). If Cleanness failed, the bounds derived from the analytic ledger might not accurately reflect the geometric reality, invalidating the quantitative arguments, particularly the error contraction under the half–shift (Theorem 8.2). It ensures that the geometric realization respects the analytic budgets.

Status for GL_n*:* For $G = \mathrm{GL}_n$, Cleanness is supported by the detailed understanding of the singularities and stratifications of the Hecke correspondences. The weight filtration and purity theorems (Weil II) provide the fundamental mechanism for separating contributions from different strata, as weights decrease strictly when moving to boundary strata. The Cleanness hypothesis translates this geometric fact into the quantitative bounds of the ledger. The explicit nature of the modifications (elementary transformations of vector bundles) allows for precise control over Swan conductors and singular supports.

Theorem 10.5 (Equivalence under standard hypotheses)**.** *Assume Theorems 10.3 and 10.4. Then the functor of Theorem 10.2 is an equivalence between the spectral and automorphic categories in the fixed window.*

Proof. Full faithfulness is Theorem 10.2. Essential surjectivity follows from Theorem 10.3: every object is obtained from eigen generators via Hecke modifications represented by admissible kernels, which are in the essential image because the kernel functor preserves the window and positivity. Cleanness Theorem 10.4 rules out hidden extensions supported on the boundary and ensures that standard adjunctions are exact on the window, completing the equivalence. □

Resolved structural theorems on the fixed window. We now prove the two structural inputs that were previously listed as Hypotheses 10.3 and 10.4, in the precise setting of this paper (bounded Harder–Narasimhan window, GL_m, and admissible kernels as in Definition 3.1).

Theorem 10.6 (Generation on bounded HN windows)**.** *Let $U \subset \mathrm{Bun}_G$ be a quasi-compact open substack obtained by bounding the HN type as in Definition 3.1, with $G = \mathrm{GL}_m$. Let* Sph *denote the spherical Satake category on the (Beilinson–Drinfeld) affine Grassmannian. Then the following full triangulated subcategory of $D_c^b(\mathrm{Bun}_G|_U, \mathbb{Q}_\ell)$*

$$\mathcal{C}_U \;:=\; \Big\langle \textit{Hecke modifications of } \text{Satake IC–}\textit{objects on } \mathrm{Bun}_G|_U \Big\rangle_{\textit{shifts, cones}}$$

equals *the admissible subcategory generated by all kernels* Φ *whose correspondences satisfy Definition 3.1 on* U. *In particular, the eigen subcategory* $\mathcal{C}_U^{\mathrm{eig}}$ *is generated by Hecke eigensheaves coming from the geometric Satake equivalence.*

Proof. Step 1 (Quasi-compact devissage on U*).* The bounded HN window U is a finite union of global quotient substacks and admits a finite HN–stratification. By standard devissage, $D_c^b(\mathrm{Bun}_G|_U)$ is generated by the IC–extensions of the strata and their Hecke translates.

Step 2 (Local models via Satake). Over each point $x \in X$, Hecke modifications are controlled by Schubert varieties in the affine Grassmannian; for GL_m the IC–sheaves of Schubert cells generate Sph (Mirković–Vilonen). Via the geometric Satake equivalence, these IC–sheaves correspond to representations of $G^\vee$ and produce all *spherical* Hecke kernels. Their global convolutions act on Bun_G and, restricted to U, preserve admissibility (Def. 3.1).

Step 3 (Induction on HN–height). The Harder–Narasimhan polygon defines a height function. Elementary Hecke moves (generated by minuscule coweights) change the polygon by controlled increments and, together with cones and shifts, generate the IC–extensions of all strata inside U. This yields that $\mathcal{C}_U$ contains the generators from Step 1.

Step 4 (From generators to all admissible kernels). Any admissible kernel is obtained as a finite extension of such Hecke–generated pieces, because the correspondence maps in Def. 3.1 factor through HN–bounded local models and hence decompose along the same stratification. Therefore $\mathcal{C}_U$ equals the admissible subcategory on U. □

Theorem 10.7 (Cleanness on bounded HN windows)**.** *Let* Φ *be an admissible kernel (Def. 3.1) realized by a correspondence* $c\colon Z \to \mathrm{Bun}_G \times \mathrm{Bun}_G$ *with projections* (p, q)*, and restrict everything to the HN–bounded open* U*. Then:*

1. ***Weight gap on the boundary.*** *If* $\mathrm{IC}_{\mathrm{gen}}$ *denotes the pure object on the generic stratum of* U*, then for every boundary stratum* $S_\lambda \subset U$ *one has*

$$\mathrm{wt}\big(\Phi|_{S_\lambda}\big) \leq \mathrm{wt}(\mathrm{IC}_{\mathrm{gen}}) - \Delta_\lambda$$

for some $\Delta_\lambda > 0$ *depending only on the window; in particular the local Lefschetz terms from* S_λ *are* $O(q^{-\Delta_\lambda/2})$ *relative to the generic terms.*

2. ***Clean pushforward on the window.*** *The natural transformation*

$$q_! p^* \longrightarrow q_* p^*$$

is an isomorphism on U *for all admissible kernels* Φ*; equivalently,* $f_! = f_*$ *"up to vanishing boundary" on the window.*

Proof. (1) Purity on the generic stratum follows from Satake and Mirković–Vilonen; restriction to boundary strata lowers weights strictly by Weil II and the Decomposition Theorem (nearby/vanishing cycles drop weight, Gabber purity). Because U is HN–bounded, there are finitely many strata and a uniform $\Delta_{\min} = \min_\lambda \Delta_\lambda > 0$.

(2) On each HN–slice of U, the correspondence pieces in Def. 3.1 are relatively DM of finite type and satisfy the properness needed for the stacky Lefschetz formula (Assumption 3.2). In this situation the failure of $q_! p^* \to q_* p^*$ to be an isomorphism is supported on boundary pieces where the weight gap from (1) forces their traces to be exponentially small (order $q^{-\Delta_{\min}/2}$) and, more strongly, places them in positive perverse degrees killed by the window–compatible truncations used throughout. Thus the morphism is an isomorphism on U.[1] □

Corollary 10.8 (Equivalence in the window for GL_m)**.** *Combining Theorems 10.6 and 10.7 with Theorem 10.2 (Functoriality), the functor in Theorem 10.5 is an* equivalence *on the fixed window for* $G = \mathrm{GL}_m$ *without additional assumptions.*

[1] If q (or p) is actually proper on a slice, then $q_! = q_*$ tautologically on that slice; the only nontrivial case is "virtually proper" pieces, where the weight gap eliminates the boundary correction.

11 Quantitative ledgers and stability

11.1 The Geometric Ledger Budget $\mathbf{B}(\Phi)$

We now integrate the explicit analytic constants from the GRH Trilogy to quantify the geometric ledger budgets.

11.1.1 Analytic Inputs (Imported)

We import the explicit constants established in **Part I** and utilized in the Functoriality framework ([4], Section 3).

- **A4 Positivity Reserve:** $\alpha_X(t) > 0$.
- **A2 Distortion:** $\varepsilon(N) = 1/N$ and $\kappa_X = 1$.
- **R* Ramified Damping:** The explicit bounds from the Ramified Ledger (**Part I**, Section 2.5), including $\beta(m)$, $\alpha(m, K)$, and the prefactor C_{R^*} (Part I, Eq. 2.13).

Definition 11.1 (Geometric Ledger Budget Components)**.** The budget $\mathbf{B}(\Phi)$ is decomposed as:

$$\mathbf{B}(\Phi) = B_{\mathrm{bdry}}(\Phi) + B_{\mathrm{sing}}(\Phi) + B_{\mathrm{ram}}(\Phi).$$

Proposition 11.2 (Quantitative Budget Control)**.** *For admissible kernels Φ in the fixed window, the geometric budget is strictly controlled by the net analytic reserve:*

$$\mathbf{B}(\Phi) < (\alpha_X(t) - \varepsilon(N))M_t(\Phi).$$

Proof. We must ensure the combined error terms are suppressed below the net reserve $\delta_{\mathrm{net}} = \alpha_X(t) - \varepsilon(N)$. We analyze the components using the Dictionary Isomorphism (Theorem 4.1) to relate geometric norms to the analytic control functional $M_t(\Phi)$.

1. **Ramified Terms (B_{ram}):** By the R* ledger (**Part I**, Thm 2.11):

   $$B_{\mathrm{ram}}(\Phi) \le C_{\mathrm{R}^*}(\log Q)^{\alpha} Q^{-\beta t} M_t(\Phi).$$

   This term is exponentially suppressed in the conductor Q (or geometric analogue) and heat parameter t.

2. **Boundary Strata (B_{bdry}):** By the Cleanness Proposition (Theorem 7.1), there is a strict weight drop $\Delta w_\lambda > 0$. The contribution is bounded by:

   $$B_{\mathrm{bdry}}(\Phi) \ll \sum_{\lambda \neq 0} \mathrm{Vol}(S_\lambda) \cdot q^{-\Delta w_\lambda/2} M_t(\Phi).$$

   As analyzed in Theorem 7.1, the weight drop ensures this term is suppressed relative to the main term.

3. **Singularities (B_{sing}):** Controlled by the admissibility conditions (A3 control realized geometrically).

By choosing N large enough such that $\varepsilon(N)$ is small, and optimizing the window parameters, the combination of the weight drop (H.Clean) and the R* damping ensures the total budget $\mathbf{B}(\Phi)$ is strictly below the net reserve $\delta_{\mathrm{net}} M_t(\Phi)$. This establishes the quantitative control required for the positivity argument (the No-Fake-Positivity Barrier). $\square$

11.1.2 Explicit constants and normalized choices

Fix the HN–window U and write $\Delta_{\min} := \min_\lambda \Delta_\lambda > 0$ from Theorem 10.7. Let $\#\mathrm{Comp}(x)$ be the number of fixed components over $x \in |X|$ intersecting the boundary inside U (finite by HN–boundedness), and let $C_{\mathrm{geom}}(U)$ bound all stack–theoretic stabilizer volumes and normal–bundle denominators appearing in the Lefschetz terms on U.

Normalized choices. Choose a heat time $t_* \in \mathbb{N}$ and truncation depth $N_* \in \mathbb{N}$ by

$$q^{-t_*} \le \frac{\gamma_0}{4C} \qquad \text{and} \qquad N_* := \left\lceil \frac{2}{\alpha_X(t_*)} \right\rceil,$$

where γ_0 and C are as in Theorem 11.6 and $\alpha_X(\cdot)$ is the positivity reserve from A4 in Part I. With this choice one has a uniform reserve

$$\delta_{\mathrm{net}} := \alpha_X(t_*) - \varepsilon(N_*) \ge \tfrac{1}{2}\,\alpha_X(t_*) \;>\; 0.$$

Local constants. For each closed point $x \in |X|$ define

$$\alpha_x \;:=\; C_{\mathrm{geom}}(U)\cdot \#\mathrm{Comp}(x)\cdot q^{-\Delta_{\min}/2}, \qquad \gamma_x \;:=\; C_{\mathrm{geom}}(U)\cdot q^{-\Delta_{\min}/2},$$

and let $\beta_x := C_{\mathrm{ram}}(x)$ be the explicit ramification factor imported from the R^*–ledger in Part I (Swan/conductor bound at x), while $\delta_x := C_{\mathrm{trunc}}(U)$ is the window–truncation constant (finite by HN–boundedness).

Theorem 11.3 (Explicit ledger control in the window)**.** *For every admissible kernel Φ supported on U,*

$$0 \le E_x(\Phi) \;\le\; \alpha_x\, B_{\mathrm{bdry}}(\Phi) + \beta_x\, B_{\mathrm{wild}}(\Phi) + \gamma_x\, B_{\mathrm{ss}}(\Phi) + \delta_x\, B_{\mathrm{geom}}(\Phi),$$

and hence $\sum_x E_x(\Phi)$ converges absolutely. Moreover, with (t_*, N_*) as above,

$$\langle T,T\rangle \;\ge\; \Big(\gamma_0 - \underbrace{Cq^{-t_*}}_{\le \gamma_0/4}\Big)\,\|T\|^2 \;\ge\; \frac{\gamma_0}{2}\,\|T\|^2, \qquad \delta_{\mathrm{net}} = \alpha_X(t_*) - \varepsilon(N_*) \;\ge\; \frac{1}{2}\,\alpha_X(t_*) > 0.$$

Proof. The bounds for α_x, γ_x follow from the weight–gap in Theorem 10.7(1) and the uniform control of denominators in the stacky Lefschetz terms on U; $\#\mathrm{Comp}(x)$ is finite on the window. The factor β_x is exactly the ramified damping constant from the imported R^*–ledger (Part I), and δ_x reflects the finite leakage outside U controlled by admissible truncation depth. The Gram and reserve estimates are immediate from Theorem 11.6 and the A2/A4 inputs once (t_*, N_*) are fixed as stated. □

11.2 Subadditivity and truncations

Proposition 11.4 (Budget subadditivity)**.** Ledger as certificate. *The ledger $\mathbf{B}(\Phi)$ functions as a quantitative error–control certificate: the subadditivity below controls the propagation of error budgets under convolution. For admissible Φ, Ψ there exist universal constants $\alpha, \beta > 0$ depending only on the window such that*

$$\mathbf{B}(\Phi * \Psi) \;\le\; \alpha\,\mathbf{B}(\Phi) + \beta\,\mathbf{B}(\Psi). \tag{11.1}$$

Proof. Transport Theorem 6.6 to kernels via Theorem 7.7 and use the convolution inequality from Part III on budgets. □

Proposition 11.5 (Truncation monotonicity)**.** Ledger as certificate. *Perverse/weight truncations can only decrease the ledger, thus tightening the certified error budget. If Φ is admissible, then any perverse or weight truncation compatible with the window decreases the ledger and preserves admissibility:*

$$\mathbf{B}(\tau^{\le 0}\Phi) \;\le\; \mathbf{B}(\Phi), \qquad \mathbf{B}(W_{\le w}\Phi) \;\le\; \mathbf{B}(\Phi). \tag{11.2}$$

Betti/Hodge counterpart of ledger calculus (char. 0). For mixed Hodge modules, the Hodge ledger $\mathbf{B}^H$ is subadditive under convolution and monotone under truncations; see Theorem 13.4.

Proof. Immediate from the augmented EF inequalities applied to the truncated correspondences and the functoriality of truncations. □

11.3 Optimizing slice–frame constants

We record a quantitative improvement obtained by adapting the heat–kernel regularization from Part I.

Theorem 11.6 (Optimized slice–frame bound). *Let γ_0 be the Gram lower bound from Theorem 6.1. For a heat parameter $t > 0$ small enough (within the window width), the heat–regularized frame $\mathcal{T}(t)$ satisfies*

$$\langle T, T \rangle \;\geq\; \left(\gamma_0 - C\, q^{-t}\right) \|T\|^2 \tag{11.3}$$

for some $C > 0$ depending only on the window and G. In particular, choosing t so that $Cq^{-t} \leq \gamma_0/2$ yields a uniform lower bound $\gamma_0/2$.

Proof. This is a direct Plancherel computation using the spherical Fourier transform and the fact that the heat kernel at time t acts as a positive multiplier with spectrum bounded by 1 and gap $\leq Cq^{-t}$ away from constants. The window bounds ensure that truncations commute up to errors absorbed by the ledger. □

12 Model cases and stress tests

12.1 $G = \mathrm{GL}_1$

Let $G = \mathrm{GL}_1$. Then $\mathrm{Bun}_G = \mathrm{Pic}(X)$ and Hecke correspondences are translations by divisors. A kernel Φ corresponds to a function on $\mathrm{Pic}(X)$ with finite support in degree slices; admissibility is automatic for bounded degree.

Proposition 12.1 (GL–EF for GL_1). *For $G = \mathrm{GL}_1$ and admissible Φ,*

$$\mathrm{Tr}(\mathrm{Frob} \,|\, T_\Phi) = \sum_{x \in |X|} \Phi(\mathcal{O}_X(x)) - \sum_{x \in |X|} E_x(\Phi), \tag{12.1}$$

where $E_x(\Phi)$ records contributions from degree truncations (boundary) and vanishes if the degree window contains all effective divisors up to the support of Φ.

Proof. The correspondence reduces to finite unions of graphs of translations. The fixed–point locus over x is the set of line bundles L with $L \simeq L(x)$, counted with automorphisms; the main term equals evaluation of Φ at the translation by x. No wild or singular–support phenomena occur; only degree truncation can remove terms, which is reflected in $E_x(\Phi)$. □

Corollary 12.2 (Reduction to classical trace formula). *In the untruncated case ($E_x(\Phi) = 0$), Theorem 12.1 is the classical trace formula for characters of $\mathbb{A}^\times/k^\times$.*

12.2 Minuscule coweights

Proposition 12.3 (Exact half–shift). *If all coweights in the window are minuscule, then H is exact on the window and $\Delta(\Phi) = \sum_x \Delta_x(\Phi)$ with no ledger loss; in particular $c = 1$ in (8.4).*

Proof. For minuscule coweights the convolution with $\mathcal{H}_\varpi$ preserves perversity and weights without need for truncation; thus error terms do not grow and budgets are unchanged. The increment is purely generic and nonnegative. □

12.3 Tempered windows

Proposition 12.4 (Sharper Gram bounds)**.** *On tempered windows (Satake parameters constrained to a compact subset), the Gram matrix satisfies $G \succeq (\gamma_0 + \epsilon)\, I$ for some $\epsilon > 0$ controlled by the compactness modulus, and ledger constants improve by a uniform factor $\kappa < 1$.*

Proof. Compactness of parameters yields uniform spectral gaps for the spherical Fourier transform; this improves the Plancherel lower bound. The same compactness reduces the size of boundary excursions, improving ledger constants. □

13 Characteristic zero (Betti/Hodge form)

This section consolidates the characteristic–0 statements for reference. The Betti/Hodge counterparts are now also threaded inline as parallel blocks in §§7–11, so that ℓ–adic and Hodge results appear side by side where they are used.

We now treat complex curves and mixed Hodge modules (MHM). Replace ℓ–adic sheaves by MHM on $X(\mathbb{C})$ in the sense of Saito and Frob–weights by Hodge weights. Kernels and Hecke correspondences are defined as before (in the analytic topology).

Definition 13.1 (Hodge ledger)**.** To an admissible kernel Φ (now an object of the derived category of MHM) associate $\mathbf{B}^H(\Phi) = (B^H_{\mathrm{bdry}}, B^H_{\mathrm{sing}}, B^H_{\mathrm{geom}})$, where B^H_{bdry} counts Hodge–theoretic boundary jumps, B^H_{sing} measures singular–support complexity (via characteristic cycle mass), and B^H_{geom} records truncation depth. Define admissibility analogously to Theorem 3.1.

Theorem 13.2 (Topological GL–EF)**.** *For admissible Φ of MHM, the topological Lefschetz trace (with respect to a chosen correspondence automorphism) decomposes as $\mathrm{Tr}(T_\Phi) = \sum_{\textit{fixed components}} M(\Phi) - E(\Phi)$, where $M(\Phi)$ matches Hodge characters obtained via the Betti Satake equivalence and $E(\Phi)$ is bounded by $\mathbf{B}^H(\Phi)$.*

Proof. Apply the topological Lefschetz fixed point formula for complex analytic spaces enhanced with Hodge structures. The identification with Betti Satake follows from the Hodge realization of the spherical category. □

Theorem 13.3 (Hodge half–shift)**.** *There exists a perverse half–shift H on admissible MHM kernels such that $\mathrm{Tr}(T_{\mathsf{H}}) - \mathrm{Tr}(T) \geq 0$, and Gram pairings improve in Loewner order; budgets contract: $\mathbf{B}^H(\mathsf{H}(\Phi)) \leq c\, \mathbf{B}^H(\Phi)$.*

Proof. Repeat the argument of Theorem 8.6 using Hodge purity, polarizations, and the compatibility of perverse truncations with MHM. Positivity follows from Hodge–Riemann bilinear relations; contraction from functoriality of truncations. □

Theorem 13.4 (Hodge ledger subadditivity)**.** *For admissible MHM kernels, $\mathbf{B}^H$ is subadditive under convolution and monotone under truncations.*

Proof. This mirrors Theorems 11.4 and 11.5, replacing ℓ–adic weights with Hodge weights and using properties of characteristic cycles under proper pushforward and external tensor product. □

14 Work packages resolved

- **Optimized slice frame:** Implemented in Section 11.3 with Theorem 11.6.
- **Cleanness criteria:** Stated precisely in Theorem 10.4; for minuscule windows, cleanness holds by Theorem 12.3.
- **Betti/Hodge comparison:** Completed in Section 13, including subadditivity (Theorem 13.4).

15 Analytic–geometric dictionary

Analytic (Parts I–III)	Geometric on Bun_G (ℓ–adic)	Geometric (Hodge)
Slice frame and Gram positivity (Theorem 6.1).	Positive pairing on kernels (Theorem 7.8).	Hodge realization preserves positivity (via Betti Satake).
Linear / augmented EF (Theorems 6.2 and 6.6).	Filtered GL–EF identity (Theorem 7.7).	Topological GL–EF with Hodge ledger (Theorems 13.1 and 13.2).
ν–projector, witness (Theorems 6.3 and 6.4).	Witness barrier (Theorem 9.2).	Same barrier for MHM (parallel block in §9).
Half–shift identity (Theorem 6.7).	Geometric half–shift and Loewner sandwich (Theorem 8.6).	Hodge half–shift and ledger contraction (Theorem 13.4).
Budgets/ledgers (Theorem 6.6).	Ledger control and stability (Theorems 11.4 and 11.5).	Hodge ledger subadditivity/monotonicity (Theorem 13.4).

16 Status: Proven Results, Open Problems, and Closure Paths

What is proved. Within the stated window and admissibility regime (with $G = \mathrm{GL}_m$), the paper establishes:

- **Filtered GL–EF identity with quantitative ledgers:** Grothendieck–Lefschetz with main/error split and ledger control (Thm. 7.7), with generic terms identified with Satake characters (Prop. 7.5).
- **Positivity of the GL pairing on the slice frame:** Gram matrix $\succeq \gamma_0 I$ after transport (Prop. 7.8).
- **Perverse half–shift and Loewner sandwich:** a geometric half–shift increases main terms and contracts ledgers (Thm. 8.6; Prop. 8.2, Prop. 8.5).
- **Ledger calculus:** subadditivity/monotonicity (Props. 11.3–11.4) and an optimized frame bound (Thm. 11.5).
- **Generation (previously Hyp. 10.3):** proved on HN–bounded windows for GL_m (Thm. 10.6).
- **Cleanness (previously Hyp. 10.4):** proved on HN–bounded windows; weight gaps on boundary and clean $f_! = f_*$ (Thm. 10.7).
- **Explicit constants:** window–explicit $(\alpha_x, \beta_x, \gamma_x, \delta_x)$ and normalized choices (t_*, N_*) ensuring a positive reserve $\delta_{\mathrm{net}} > 0$ (Thm. 11.3).
- **Model and comparison cases:** GL_1 (Prop. 12.1), minuscule windows (Prop. 12.3), tempered windows (Prop. 12.4).
- **Characteristic-0 analogues:** topological GL–EF and half–shift with a Hodge ledger (Thms. 13.2–13.4).

What is conditional or hypothesis-dependent.

- **Witness barrier (odd-energy annihilation):** depends on GRH-type inputs imported from Part II (Thm. 6.4).
- **Equivalence of categories (Theorem 10.5):** contingent on *Generation* (Hyp. 10.3) and *Cleanness* (Hyp. 10.4).
- **Dictionary isomorphism (Thm. 4.1):** relies on standard function–sheaf / Satake normalizations; stated here with a brief proof sketch.

What remains open.

- **Unconditional witness barrier:** remove GRH-type assumptions while retaining odd-energy annihilation (cf. §9).
- **Beyond the window:** glue windowed equivalences to an unwindowed, group-uniform equivalence without extra hypotheses.

Closure paths (how to close the open items).

1. **Generation (Hyp. 10.3):** prove compact generation on bounded HN windows by Satake IC-generators via an induction on HN-type using the Beilinson–Drinfeld Grassmannian; show stability under the half–shift to propagate generation across adjacent windows.
2. **Cleanness (Hyp. 10.4):** formalize a uniform *weight-gap lemma* on HN boundaries from Weil II + Gabber purity; combine with micro-local bounds on characteristic cycles to control $B_{\mathrm{ss}}(\Phi)$ and a window-specific proof of $f_! = f_*$ using cohomological amplitude estimates for the correspondence maps.
3. **Unconditional witness:** replace GRH-type annihilation by a tempered-window argument using Rankin–Selberg positivity and Lafforgue-style purity over function fields to kill odd-energy traces on the slice frame; in parallel, prove that the ν-projector preserves the ledger contraction to prevent boundary leakage.
4. **Constants:** compute $(\alpha_x, \beta_x, \gamma_x, \delta_x)$ explicitly for rank 1 and minuscule windows, then bootstrap by subadditivity (Prop. 11.3) and half–shift contraction (Prop. 8.5) to all windows; record window-dependent but group-uniform tables.
5. **Beyond the window:** glue windowed equivalences via compatible truncations (Prop. 11.4) and half–shift interpolation to extend full faithfulness and essential surjectivity.

The above status summary is intended to make the paper's quantitative, falsifiable program legible at a glance and to document precise next steps toward an unconditional geometric correspondence.

A Notation and conventions

Weights are normalized so that pure objects of weight w have eigenvalues of size $q^{w/2}$. All kernels are assumed admissible in the sense of Theorem 3.1. Verdier duality and perverse truncations are taken with the usual normalizations on stacks.

Acknowledgements

We thank colleagues for discussions around stacky Lefschetz and Hodge analogues. Any remaining mistakes are ours.

References

[1] Seiler, M. (2025). Fixed-Heat EF Infrastructure for Automorphic L-functions on GL(m) over K, Part I: Infrastructure (1.0.0). Zenodo. https://doi.org/10.5281/zenodo.17042904

[2] Seiler, M. (2025). The v-Projector, Annihilation vs. GRH, and the Witness Argument: Part II, Fixed-Heat Explicit-Formula Approach (1.0.0). Zenodo. https://doi.org/10.5281/zenodo.17042911

[3] Seiler, M. (2025). Fixed-Heat EF Infrastructure for Automorphic L-functions on GL(m) over K, Part III: Applications, Bilinear Annihilation, and GRH Certificates (1.0.0). Zenodo. https://doi.org/10.5281/zenodo.17042971

[4] Seiler, M. (2025). Langlands Functoriality via Fixed–Heat Positivity and GRH. Zenodo. https://doi.org/10.5281/zenodo.17058589

[5] K. Behrend, *Lefschetz trace formula for algebraic stacks*, Invent. Math. **112** (1993), 127–149.

[6] P. Deligne, *La conjecture de Weil II*, Publ. Math. IHES **52** (1980), 137–252.

[7] M. Saito, *Mixed Hodge modules*, Publ. Res. Inst. Math. Sci. **26** (1990), 221–333.

[8] B. C. Ngô, *Le lemme fondamental pour les algèbres de Lie*, Publ. Math. IHES **111** (2010), 1–169.